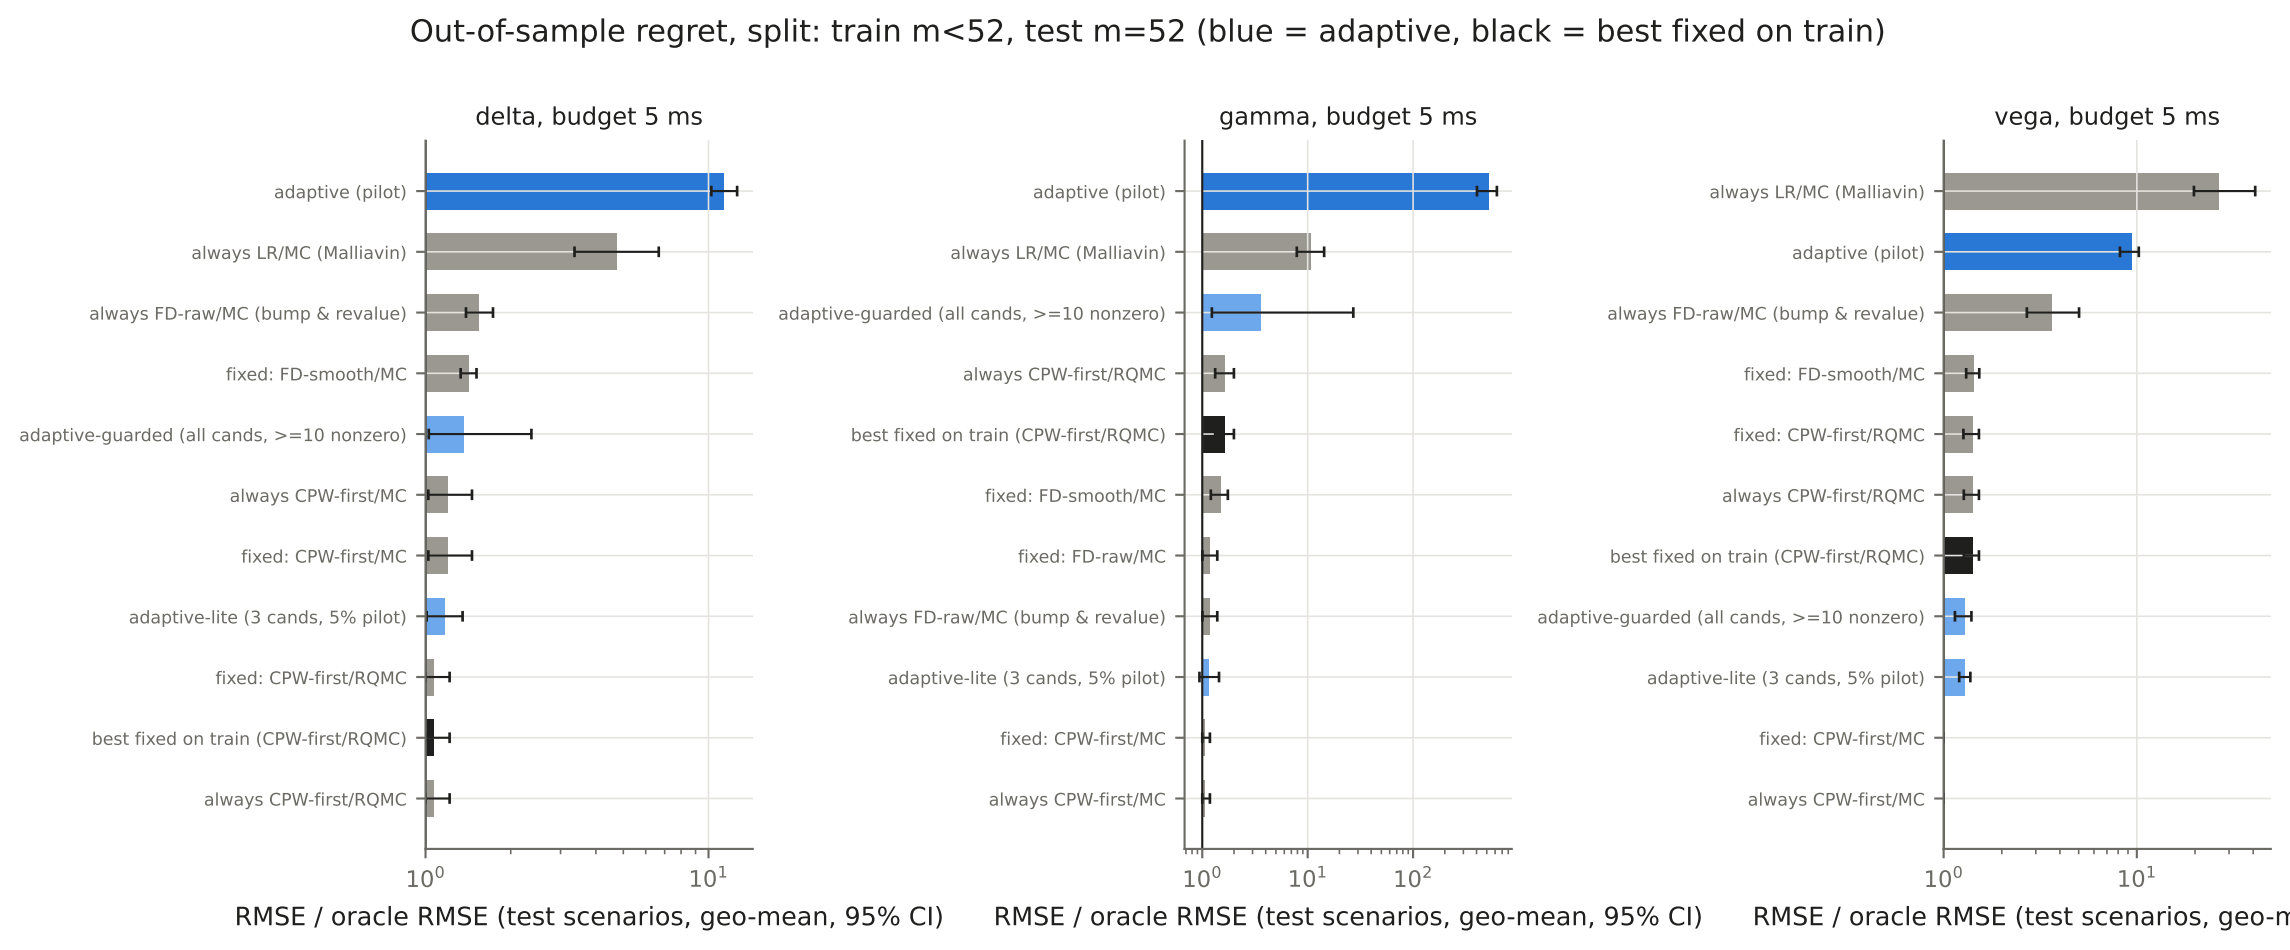
Out-of-sample regret, split: train m<52, test m=52 (blue = adaptive, black = best fixed on train)
delta, budget 5 ms
adaptive (pilot)
always LR/MC (Malliavin)
always FD-raw/MC (bump & revalue)
fixed: FD-smooth/MC
adaptive-guarded (all cands, >=10 nonzero)
always CPW-first/MC
fixed: CPW-first/MC
adaptive-lite (3 cands, 5% pilot)
fixed: CPW-first/RQMC
best fixed on train (CPW-first/RQMC)
always CPW-first/RQMC
10^0
10^1
RMSE / oracle RMSE (test scenarios, geo-mean, 95% CI)
gamma, budget 5 ms
adaptive (pilot)
always LR/MC (Malliavin)
adaptive-guarded (all cands, >=10 nonzero)
always CPW-first/RQMC
best fixed on train (CPW-first/RQMC)
fixed: FD-smooth/MC
fixed: FD-raw/MC
always FD-raw/MC (bump & revalue)
adaptive-lite (3 cands, 5% pilot)
fixed: CPW-first/MC
always CPW-first/MC
10^0
10^1
10^2
RMSE / oracle RMSE (test scenarios, geo-mean, 95% CI)
vega, budget 5 ms
always LR/MC (Malliavin)
adaptive (pilot)
always FD-raw/MC (bump & revalue)
fixed: FD-smooth/MC
fixed: CPW-first/RQMC
always CPW-first/RQMC
best fixed on train (CPW-first/RQMC)
adaptive-guarded (all cands, >=10 nonzero)
adaptive-lite (3 cands, 5% pilot)
fixed: CPW-first/MC
always CPW-first/MC
10^0
10^1
RMSE / oracle RMSE (test scenarios, geo-m