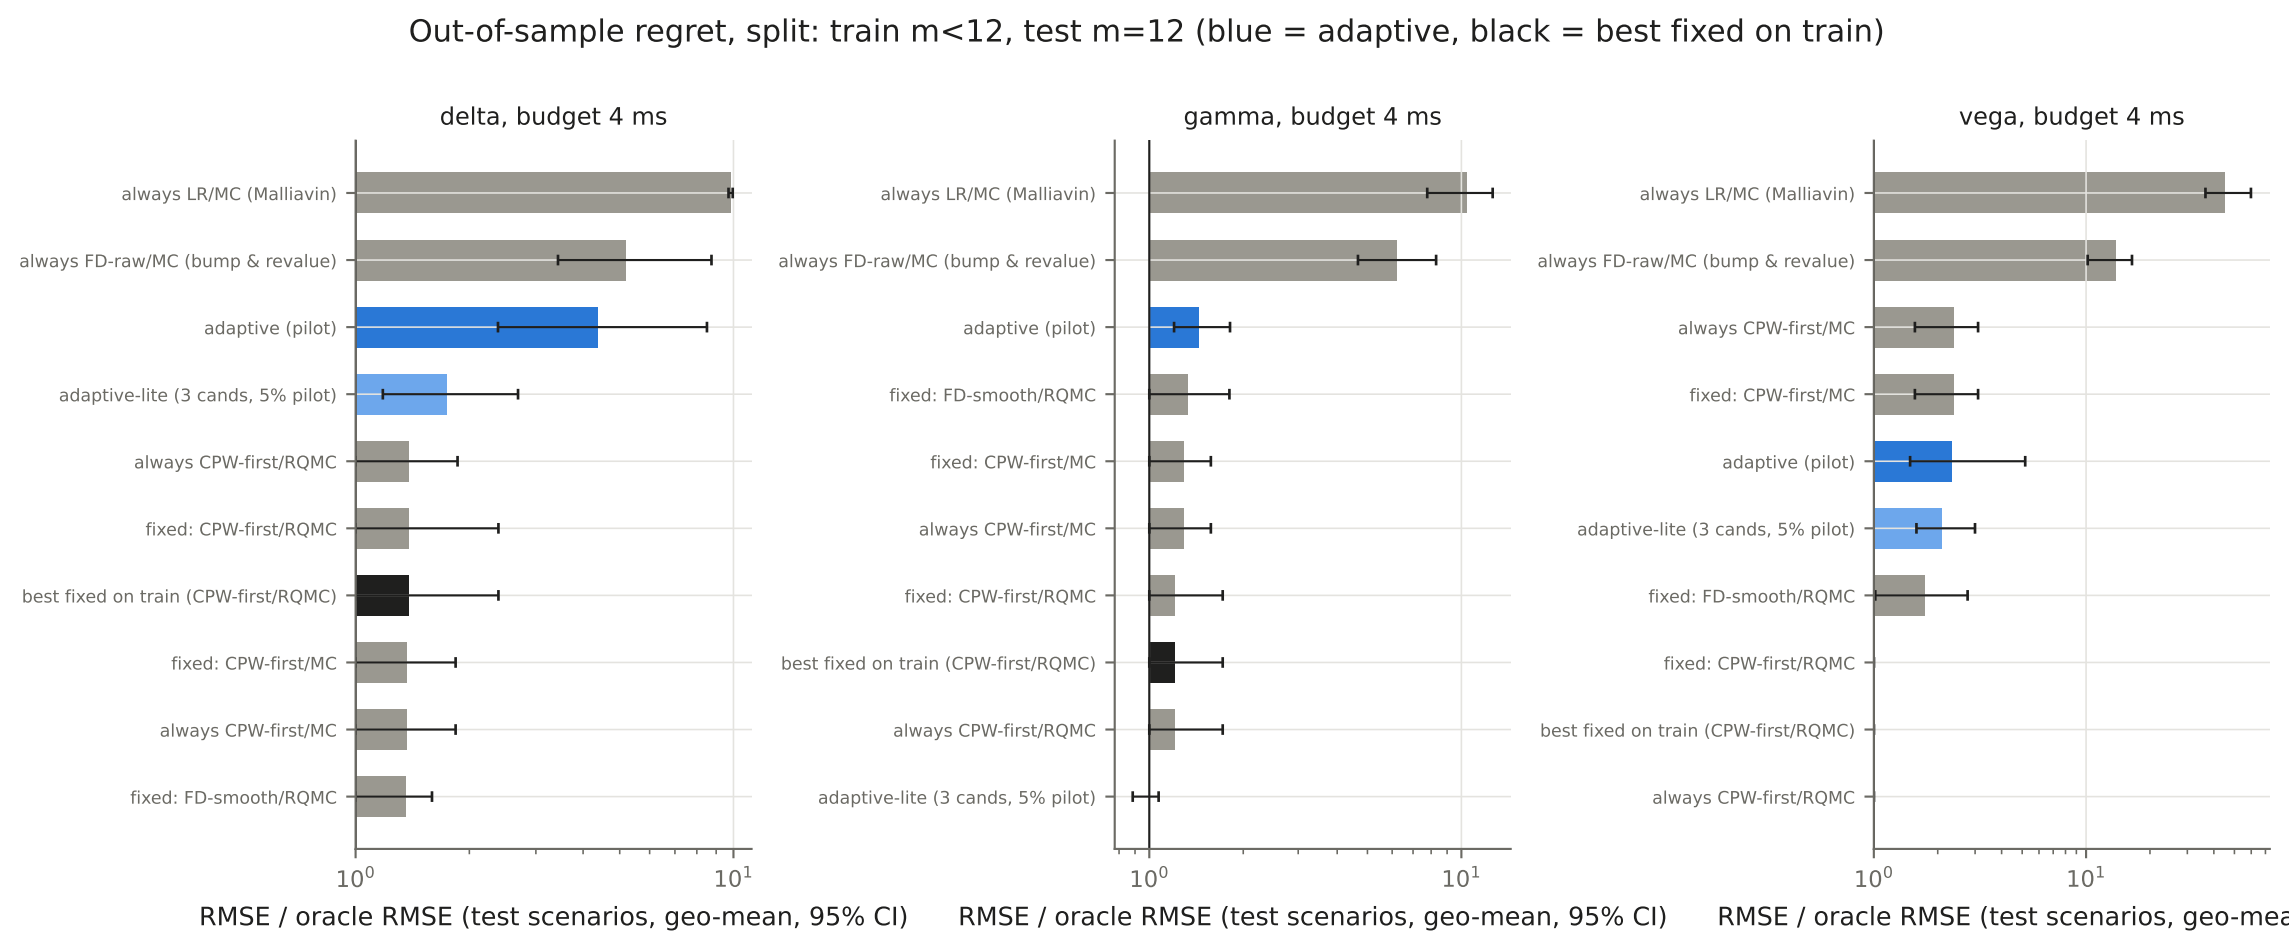
Out-of-sample regret, split: train m<12, test m=12 (blue = adaptive, black = best fixed on train)
delta, budget 4 ms
gamma, budget 4 ms
vega, budget 4 ms
always LR/MC (Malliavin)
always FD-raw/MC (bump & revalue)
adaptive (pilot)
adaptive-lite (3 cands, 5% pilot)
always CPW-first/RQMC
fixed: CPW-first/RQMC
best fixed on train (CPW-first/RQMC)
fixed: CPW-first/MC
always CPW-first/MC
fixed: FD-smooth/RQMC
10^0
10^1
RMSE / oracle RMSE (test scenarios, geo-mean, 95% CI)