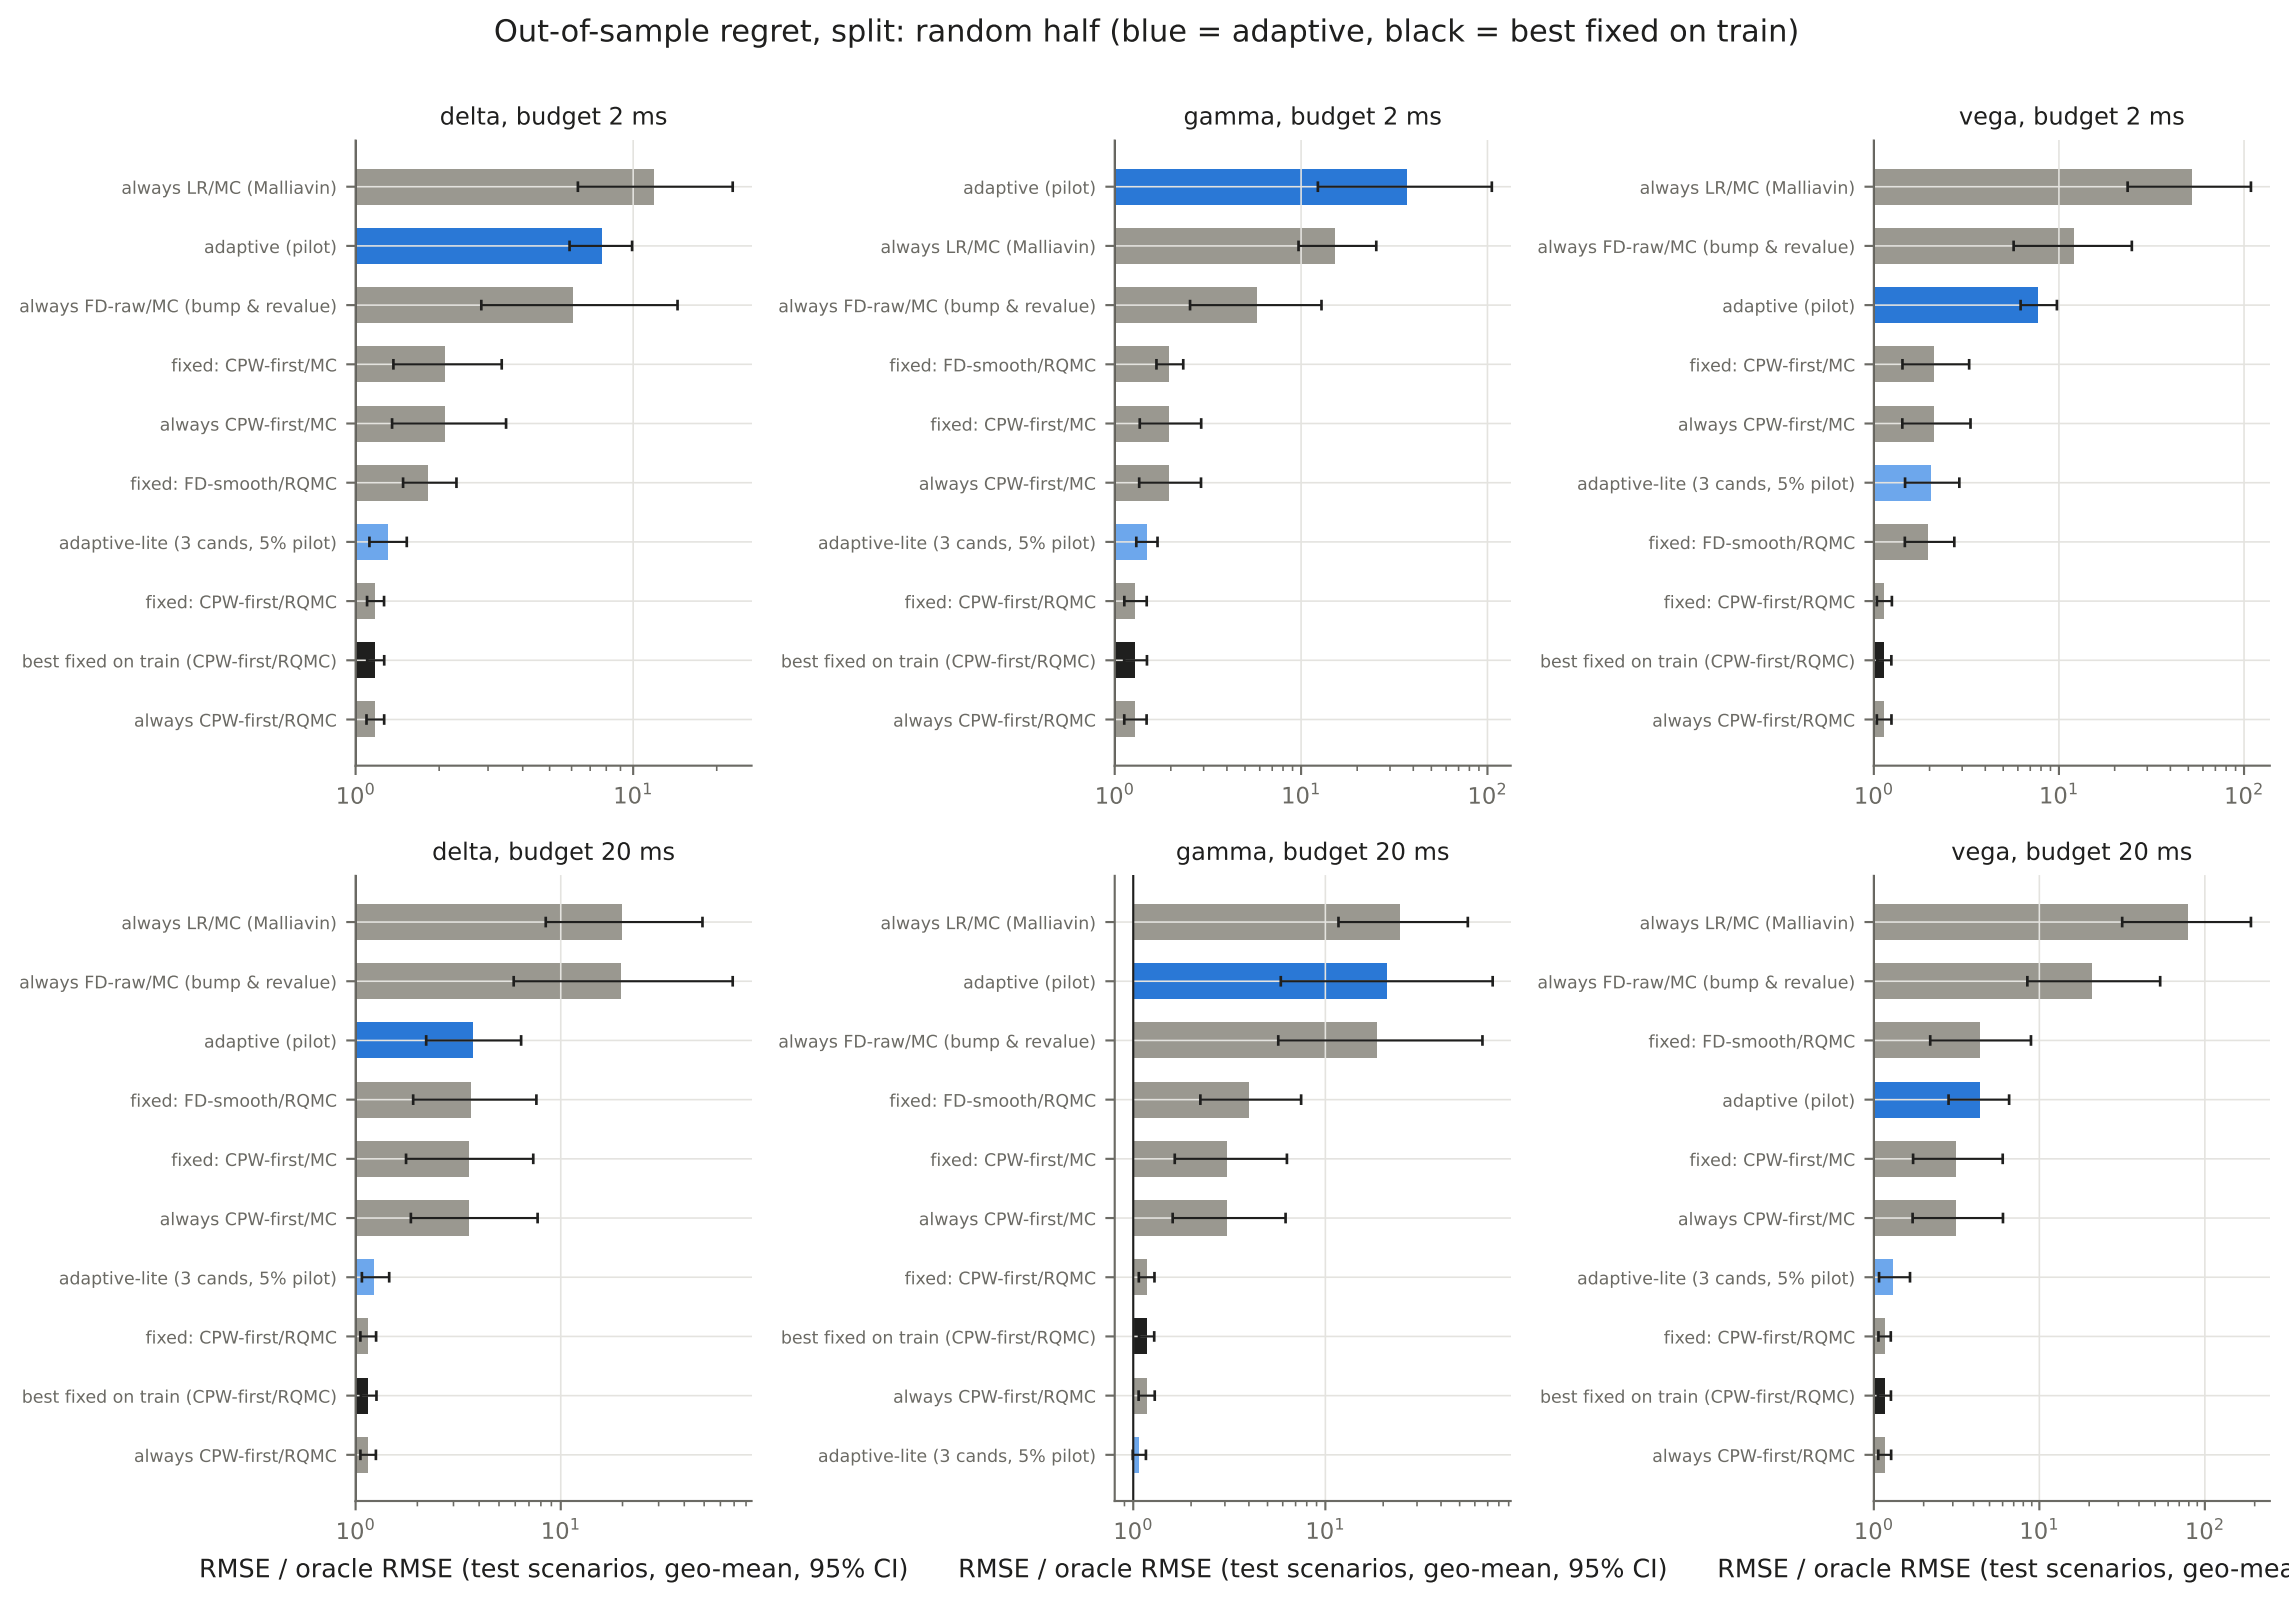
Out-of-sample regret, split: random half (blue = adaptive, black = best fixed on train)
delta, budget 2 ms
gamma, budget 2 ms
vega, budget 2 ms
delta, budget 20 ms
gamma, budget 20 ms
vega, budget 20 ms
always LR/MC (Malliavin)
adaptive (pilot)
always FD-raw/MC (bump & revalue)
fixed: CPW-first/MC
always CPW-first/MC
fixed: FD-smooth/RQMC
adaptive-lite (3 cands, 5% pilot)
fixed: CPW-first/RQMC
best fixed on train (CPW-first/RQMC)
always CPW-first/RQMC
10^0
10^1
10^2
RMSE / oracle RMSE (test scenarios, geo-mean, 95% CI)