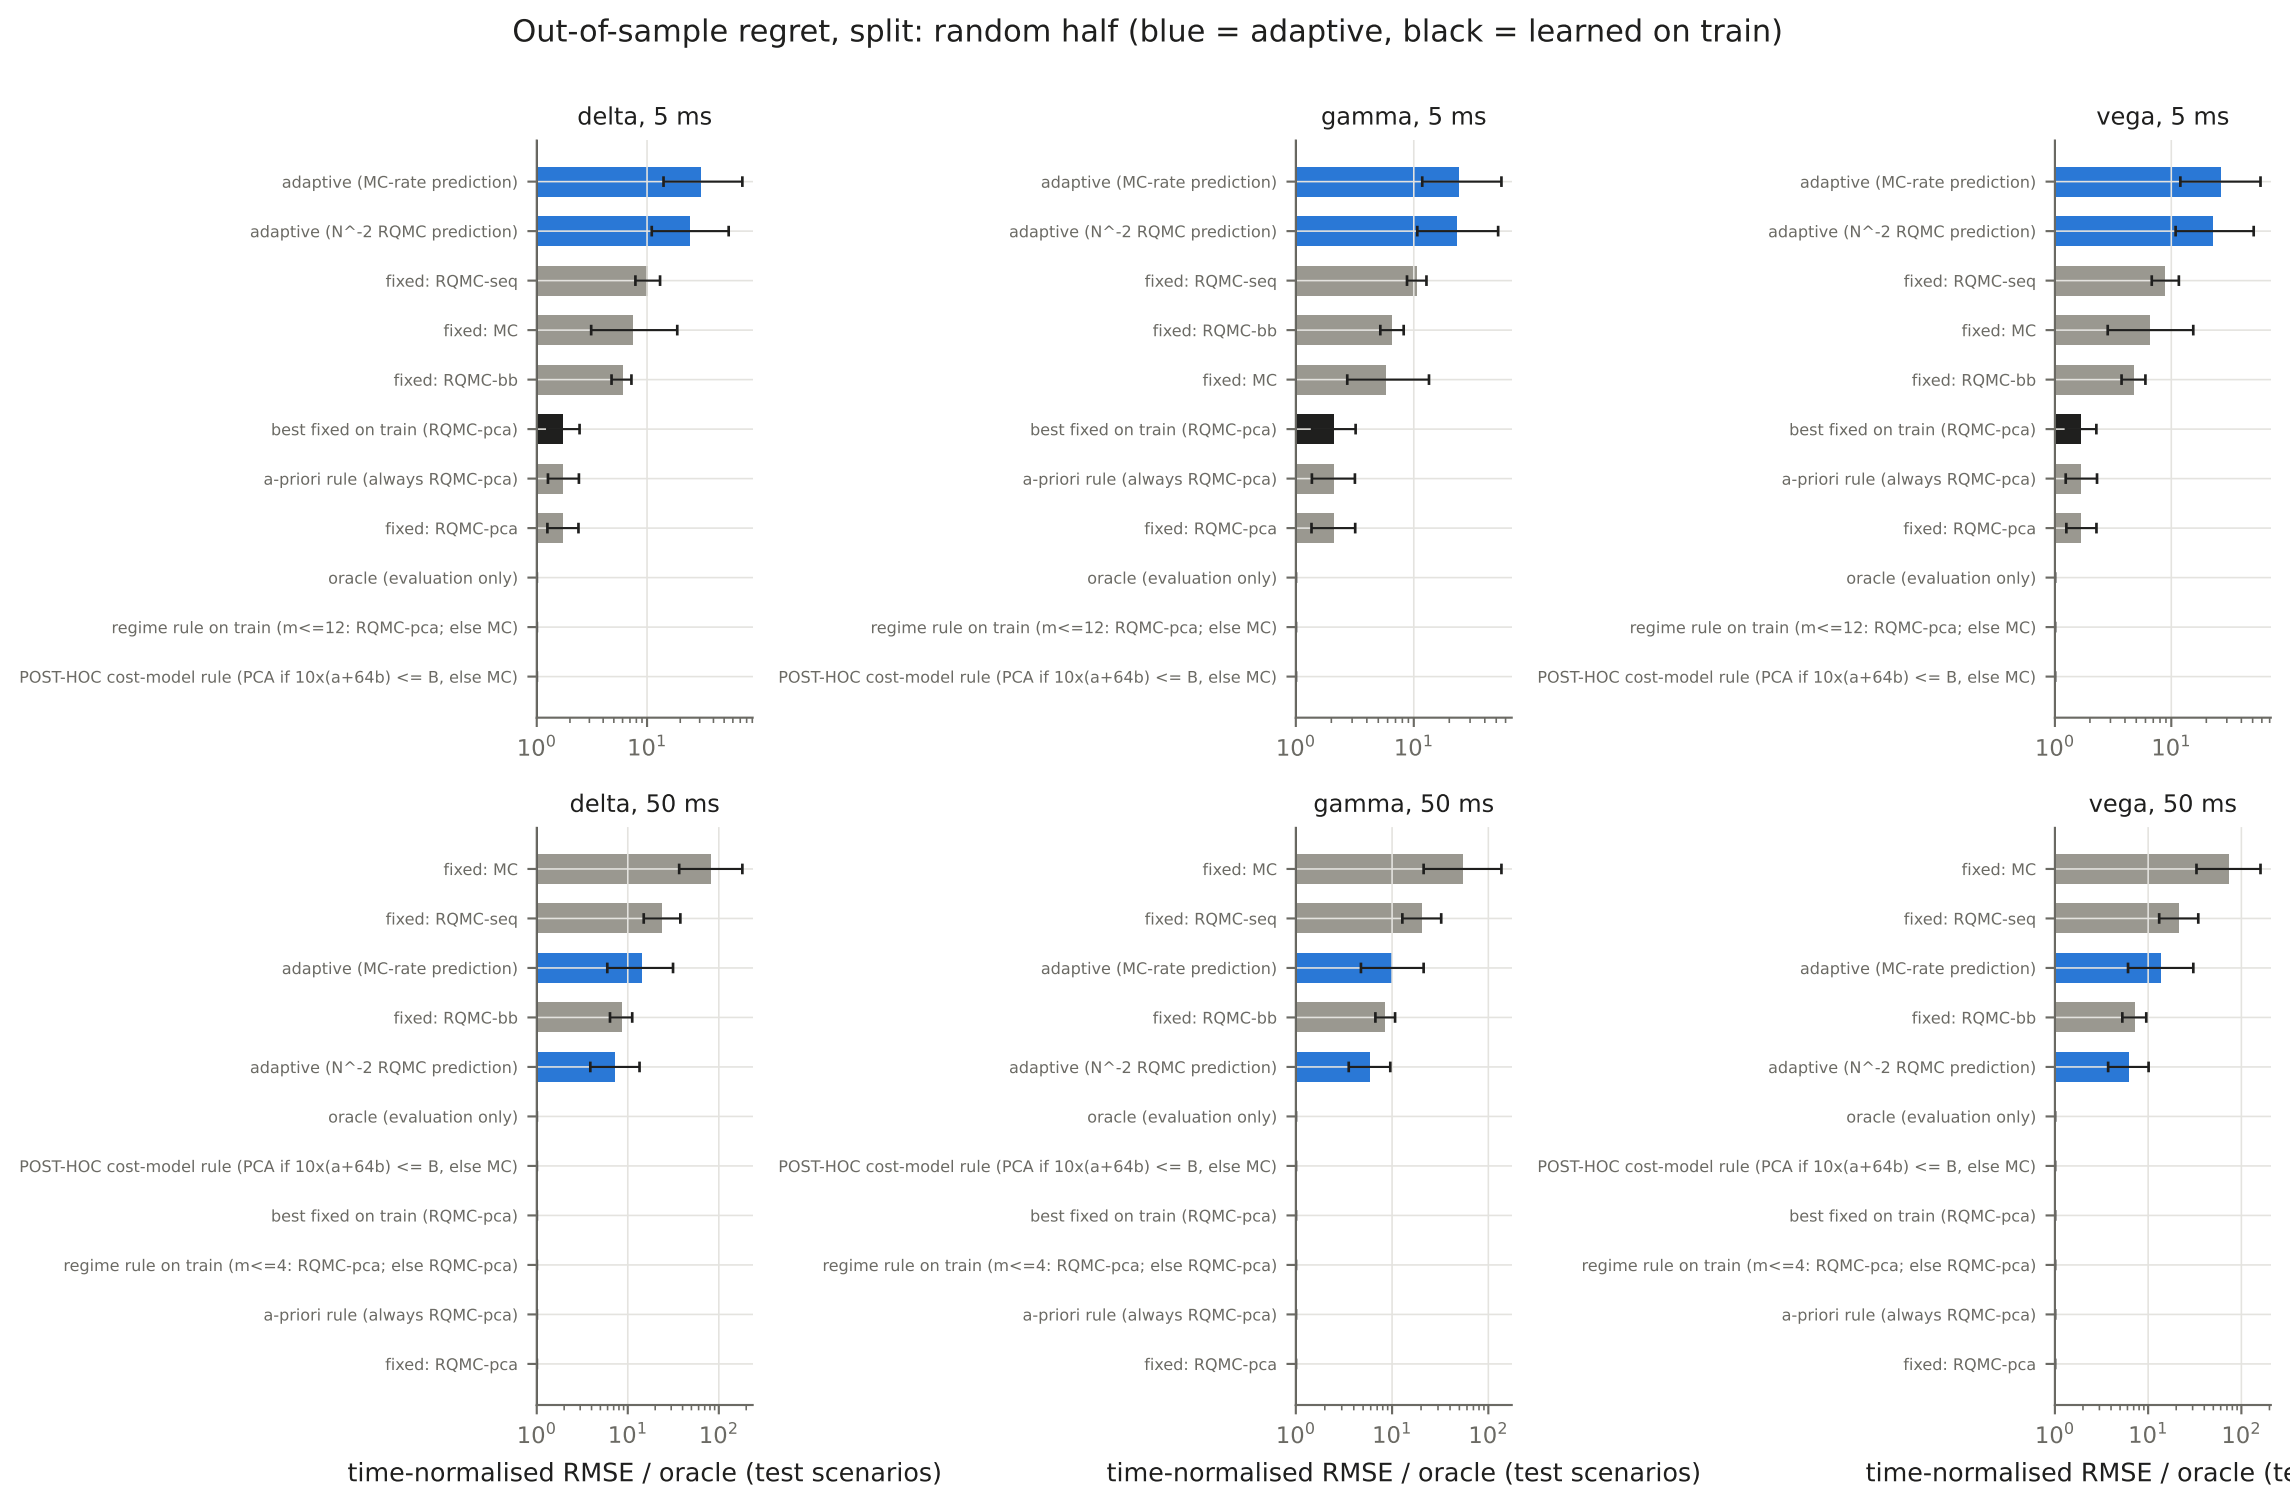
Out-of-sample regret, split: random half (blue = adaptive, black = learned on train)
delta, 5 ms
gamma, 5 ms
vega, 5 ms
adaptive (MC-rate prediction)
adaptive (N^-2 RQMC prediction)
fixed: RQMC-seq
fixed: MC
fixed: RQMC-bb
best fixed on train (RQMC-pca)
a-priori rule (always RQMC-pca)
fixed: RQMC-pca
oracle (evaluation only)
regime rule on train (m<=12: RQMC-pca; else MC)
POST-HOC cost-model rule (PCA if 10x(a+64b) <= B, else MC)
10^0
10^1
delta, 50 ms
gamma, 50 ms
vega, 50 ms
fixed: MC
fixed: RQMC-seq
adaptive (MC-rate prediction)
fixed: RQMC-bb
adaptive (N^-2 RQMC prediction)
oracle (evaluation only)
POST-HOC cost-model rule (PCA if 10x(a+64b) <= B, else MC)
best fixed on train (RQMC-pca)
regime rule on train (m<=4: RQMC-pca; else RQMC-pca)
a-priori rule (always RQMC-pca)
fixed: RQMC-pca
10^0
10^1
10^2
time-normalised RMSE / oracle (test scenarios)
time-normalised RMSE / oracle (test scenarios)
time-normalised RMSE / oracle (te